

Layer-skipping connections facilitate training of layered networks using equilibrium propagation.

Jimmy Gammell
Sae Woo Nam
Adam N. McCaughan

ABSTRACT

Equilibrium propagation is a learning framework for energy-based networks that can be implemented by neurons that perform only one type of computation in the prediction and correction phases of training, and that computes parameter corrections for a given neuron using only the activation values of directly-connected neurons. This makes it an appealing candidate for implementation in neuromorphic analog hardware, and marks a step forward in the search for a biologically-plausible implementation of deep learning. However, in previous implementations of equilibrium propagation, layered networks suffered from a vanishing gradient problem that has not yet been solved in a simple, biologically-plausible way. In this paper, we demonstrate that the vanishing gradient problem can be counteracted by replacing some of a layered network's connections with random layer-skipping connections in a manner inspired by small-world networks. This approach could be conveniently implemented in neuromorphic analog hardware and is biologically-plausible.

CCS CONCEPTS

• Computing methodologies → Bio-inspired approaches; Neural networks.

KEYWORDS

deep learning, neural network, biologically-motivated, equilibrium propagation, vanishing gradient, small-world, neuromorphic hardware

ACM Reference Format:

Jimmy Gammell, Sae Woo Nam, and Adam N. McCaughan. 2020. Layer-skipping connections facilitate training of layered networks using equilibrium propagation.. In *Proceedings of International Conference on Neuromorphic Systems (ICONS 2020)*. ACM, New York, NY, USA, 7 pages. <https://doi.org/10.1145/nnnnnnnn.nnnnnnnn>

1 INTRODUCTION

Equilibrium propagation [18] is a learning framework for energy-based networks such as the continuous Hopfield network [8]. It is

appealing relative to backpropagation because it is more biologically-plausible, and as a side-effect could be implemented more-easily in neuromorphic analog hardware.

A major reason backpropagation is not biologically-plausible is that each neuron would need two distinct mechanisms for information transmission: one in the forward propagation phase to transmit its activation to shallower neurons, and another in the backpropagation phase to transmit error-correction information to deeper neurons [2]. While this is easy on a digital computer that oversees and manipulates the entire network, it would be cumbersome in hardware (biological or otherwise) consisting of many simple, independent computational nodes with limited ability to share information. In contrast, equilibrium propagation consists of a free phase (comparable to forward-propagation) and a weakly-clamped phase (comparable to backwards-propagation) that can be carried out using only the activations of neighboring neurons, necessitating only one mechanism for information transmission [18].

In a similar vein, in backpropagation each neuron would need to perform two distinct computations: computation of an activation using deeper neurons' activations in the forward-propagation phase, and computation of an error term using shallower neurons' activation terms in the backwards-propagation phase. In equilibrium propagation the only difference between the phases is that in the weakly-clamped phase the activations of a subset of neurons are weakly-clamped towards a target output; the dynamics are otherwise identical [18].

Implementation of equilibrium propagation in [18] was hindered by a vanishing gradient problem whereby networks with as few as 3 hidden layers trained slowly on MNIST [12] - a serious issue given that network depth is critical to performance on difficult datasets [21, 22] and that convergence to a low error rate on MNIST is a low bar to meet. The problem was overcome in [18] by independently tuning a unique learning rate for each layer in the network; however, this approach is unappealing because (1) it introduces additional hyperparameters to tune, (2) it would be inconvenient to implement in neuromorphic analog hardware, and (3) it has not been observed in biological systems.

The purpose of this paper is to demonstrate that in this context the vanishing gradient problem can instead be solved by randomly replacing some of a layered network's connections with layer-skipping connections.¹ Through this modification we have achieved 0% training error (out of 50,000 examples) and $\leq 2.5\%$ test error (out of 10,000 examples) on MNIST using a network with three

Permission to make digital or hard copies of all or part of this work for personal or classroom use is granted without fee provided that copies are not made or distributed for profit or commercial advantage and that copies bear this notice and the full citation on the first page. Copyrights for components of this work owned by others than ACM must be honored. Abstracting with credit is permitted. To copy otherwise, or republish, to post on servers or to redistribute to lists, requires prior specific permission and/or a fee. Request permissions from permissions@acm.org.

ICONS 2020, July 2020, Chicago, Illinois

© 2020 Association for Computing Machinery.

ACM ISBN 978-x-xxxx-xxxx-x/YY/MM...\$15.00

<https://doi.org/10.1145/nnnnnnnn.nnnnnnnn>

¹This modification was inspired by small-world topology [24]; however, we have not observed a strong correlation between network performance and common metrics of small-worldness (characteristic path length, clustering coefficient, small-world coefficient [9]).

hidden layers and no regularization term in its cost function. These error rates are similar to those of other biologically-motivated networks [1] and are roughly the same as those of the layered network with unique, manually-tuned learning rates in [18], although they required around 25% more epochs to achieve. Our method adds only one additional hyperparameter - the number of connections to randomly replace - while removing $L - 2$ hyperparameters, where L is the number of layers in the network,² and could be implemented with relative ease in any system with configurable connectivity. Layer-skipping connections have been observed in biological brains [3], so the approach is biologically-plausible. Similar techniques have seen success in convolutional [7, 23] and multilayer feedforward [11, 25] networks. Our findings in this paper suggest that layer-skipping connections are effective-enough to be appealing in contexts where simplicity and biological plausibility are important.

2 BACKGROUND AND THEORY

2.1 Equilibrium propagation

Equilibrium propagation [18] is a learning framework for energy-based networks that trains their parameters by approximating gradient descent on some arbitrary cost function. It is applicable to any network with dynamics characterized by evolution to a fixed point of an associated energy function; in this and in [18], it is implemented on a continuous Hopfield network [8]. For a variety of reasons outlined in [2], backpropagation is not biologically-plausible. One of the major reasons is that to correct the parameters of a given neuron, backpropagation requires precise information about the activations and nonlinearities of all neurons in the corresponding feedforward path; equilibrium propagation avoids this problem by correcting a neuron's parameters using only the activations of neurons to which it directly connects. Another major reason backpropagation is not biologically-plausible is that it requires neurons to perform distinct computations (implemented using distinct circuitry) in the forward and backward propagation phases of training; in contrast, equilibrium propagation requires only one behavior of neurons in both the prediction (free) and correction (weakly-clamped) phases of training. In addition to enhancing its biological plausibility, these traits make equilibrium propagation appealing as a framework to implement on neuromorphic analog hardware because they would limit the complexity required of neurons and the amount of infrastructure needed to operate and train them.

2.2 Vanishing gradient problem

It has been observed [18] that energy-based networks with a layered topology trained using equilibrium propagation suffer from a vanishing gradient problem whereby the rate of training of a weight decreases approximately exponentially with the number of layers between that weight and the output layer. This is problematic because it has been observed to increase the time needed to train a network, and gradients spanning several orders of magnitude could be hard to represent in analog implementations of the framework with limited bit depth. The vanishing gradient problem is familiar

in the context of conventional networks trained through backpropagation, where the composition of many layers can impede training by amplifying or attenuating the effect of changes to deep neurons as these changes propagate to the output layer, leading to a vanishing or exploding gradient. For those networks the problem can be effectively addressed by initializing weights to make activation variances and backpropagated gradient variances approximately uniform with respect to depth as described in [6], and by using activation functions with derivatives that do not cause neurons to saturate [19], such as rectified linear units $\rho(x) = \text{Max}\{0, x\}$ or hardened sigmoids (equation ??). Batch normalization [10] is another effective technique for promoting uniform training in a network; inputs to each layer are whitened (the input distribution \mathbf{X} to each layer is transformed to have $E[\mathbf{X}] = \mathbf{0}$ and $\text{Var}[\mathbf{X}] = \mathbf{1}$) during each batch to prevent changes to deep parameters from changing the input distribution to shallower layers and requiring adjustment to compensate.

The weight initialization scheme from [6] and the use of hardened sigmoid activation functions do not solve the vanishing gradient problem in the context of equilibrium propagation. Batch normalization has not been tried in this context, but is not biologically-plausible. This suggests that the cause of the vanishing gradient problem here is different than in a conventional network. The situation in energy-based networks is more-complicated because there is no straightforward relationship between the activations of two given neurons; instead, neurons evolve as a coupled dynamical system.

In [18] the vanishing gradient problem was addressed by using unique learning rates for each layer, chosen to make the magnitudes of corrections to neurons' parameters uniform regardless of depth in the network. While this method is effective, it is unappealing because it seems unlikely to take place in a biological brain, it would add complexity to an analog implementation of the framework, and it introduces more hyperparameters that must be tuned to train a network. Our topology effectively addresses the problem without these unappealing characteristics.

3 IMPLEMENTATION

We implemented³ the equilibrium propagation learning framework [18] using the Pytorch library, and verified that we could recreate the experiments ran in [18]. We follow the procedure described in section 2.1; we use a hardened sigmoid activation function (equation ??) and a squared-error cost function with no regularization term (equation ??). Testing was done on the MNIST dataset [12] with input-output pairs grouped into batches of 20; training was done on the 50,000 training pairs and testing on the 10,000 validation pairs.

We use two performance-boosting techniques described in [18]. We randomize the sign of β at the outset of each batch, which provides a regularization effect (note that parameter corrections are divided by β , so are changed to minimize cost regardless of its sign). We use persistent particles: after the first epoch, before training on a batch the network's state is initialized to its free equilibrium state

²Our current implementation also requires manual tuning of the initial weights of intralayer and layer-skipping connections, but we have found that networks are not very sensitive to these initializations and that performance is good as long as they are in the same ballpark as the rest of the network's initializations.

³https://github.com/jgammell/Equilibrium_Propagation_mobile.git




Figure 1: Topology of the layered network tested in [18]. All pairs of neurons in adjacent layers are connected, and there are no additional connections. All connections are bidirectional. The learning rate for weights is reduced by a factor of 4 each time distance from the output decreases by one layer, to compensate for the vanishing gradient problem.




Figure 2: Changes to the layered topology to compensate for the vanishing gradient problem while using a single learning rate for all weights. Red dotted lines denote connections that have been removed and black dotted lines denote their replacements. Green solid lines denote added intralayer connections. In this illustration connections have been replaced by a random layer-skipping connection with probability $p \approx 8\%$. All connections are bidirectional.

from the last time it trained on the same batch. This reduces computational resources needed for training by reducing the number of iterations of equation ?? needed to reach equilibrium, and would be unnecessary in an analog implementation; note that training and test error rates during early epochs are artificially inflated as a result of this technique.

3.1 Layered topology with unique learning rates

We recreated the 5-layer network used in [18]. It has a standard layered topology shown in 1.

The network consists of an input layer with 784 neurons, 3 hidden layers with 500 neurons each, and an output layer with 10 neurons. There are connections between every pair of neurons in adjacent layers, no connections between neurons in the same layer, and no layer-skipping connections.

The weight matrix was initialized using the Glorot-Bengio initialization scheme [6]: a weight connecting layer j with n_j neurons to layer $j + 1$ with n_{j+1} neurons is drawn from a uniform distribution

$$U\left[-\frac{\sqrt{6}}{\sqrt{n_j + n_{j+1}}}, \frac{\sqrt{6}}{\sqrt{n_j + n_{j+1}}}\right]. \quad (1)$$

For effective training it is necessary to use unique learning rates to train the weights connecting each pair of layers. We use the same learning rates that were used for experiments in [18]: .128, .032, .008, .002, in order of the deepest to the shallowest layer; to counter the vanishing gradient problem the learning rate is reduced by a factor of 4 each time the distance from the output layer decreases by one layer. Each neuron has a bias term that is trained using the learning rate corresponding to the weights of the neuron's input connections, e.g. the biases of the first hidden layer are trained with a learning rate of .128.

3.2 Layered topology with single learning rate

To provide a point of reference, we ran tests on a network that is identical to that in section 3.1 except with a single learning rate of .02 across the entire network. It performs poorly due to the vanishing gradient problem.

3.3 Our topology

Algorithm 1: Algorithm to produce our topology

Input: A network N with layered topology
 // described in section 3.2
Input: An integer n
 // number of connections to replace
Output: A network with our modified topology
for hidden layer $l \in N$ **do**
 | Add edge between each pair of neurons in l
for $i \leftarrow 1$ **to** n **do**
 | Randomly select connection $c \in N$;
 | Add connection between random unconnected pair of neurons $\in N$;
 | // Do not allow self connections
 | // Do not allow connections between two input neurons or between two output neurons
 | Remove c ;
return N

To generate a network with our topology, we use algorithm 1. The topology is illustrated in figure 2. For $p \lesssim .2$, algorithm 1 is roughly equivalent to algorithm ?? with $p = 1 - (\frac{N_o - 1}{N_o})^n$, where

N_o is the number of connections in the network⁴; to contextualize the number of replaced connections, we will henceforth describe networks with our topology in terms of p and not n . We have also tried adding new connections without removing existing ones, and observed roughly the same performance; in this paper we replace connections to limit the number of parameters added to the network. We do not add connections within the input or output layers and we do not allow neurons to connect to themselves. We have seen good results with $p \approx 8\%$.

This topology allows us to train the network with a uniform learning rate of .02 across the entire network. Weights of connections between pairs of neurons in adjacent layers are still initialized based on equation 1. Weights of intralayer connections and layer-skipping connections are drawn from the uniform distribution $U[-.05, .05]$, where the value .05 was chosen empirically to optimize network performance.

3.4 Tracking the training rates of individual pairs of layers

We find it informative to track the root-mean-square magnitude of corrections to weights in individual layers; if $\Delta w_{ij}^l(b)$ denotes the correction to the connection weight between the i^{th} neuron in layer l and the j^{th} neuron in layer $l + 1$ in response to batch b and N^l denotes the number of neurons in layer l , we define

$$\Delta w^l(b) := \sqrt{\frac{\sum_{i=1}^{N^l} \sum_{j=1}^{N^{l+1}} (\Delta w_{ij}^l(b))^2}{N^l N^{l+1}}} \quad (2)$$

as a metric of the extent to which the weights between layers l and $l + 1$ trained in response to batch b . We ignore the training of the weights of intralayer and layer-skipping connections.

This measurement is volatile, so for clarity we plot the average of $\Delta w^l(b')$ for b' in a neighborhood of batch b : for some n , we define

$$\Delta \hat{w}^l(b) := \frac{1}{2n+1} \sum_{b'=b-n}^{b+n} \Delta w^l(b'). \quad (3)$$

Comparing the traces $\Delta \hat{w}^l(b)$ vs. b for different layers reveals the extent of the vanishing gradient problem.

4 RESULTS

We compared the three network topologies described in sections 3.1, 3.2, and 3.3 on the MNIST dataset [12]. The only hyperparameter that differed between networks was the learning rate(s); the layered network with unique learning rates used rates of .128, .032, .008, and .002 for layers from input, and the layered network with one learning rate and the network with our topology used a learning rate of .02. All networks used $\epsilon = .5$, $\beta = 1.0$, $N_{free} = 500$, $N_{weakly-clamped} = 8$ and were trained for 250 epochs with 50,000 training examples, 10,000 test examples and a batch size of 20. To generate our topology we used algorithm 1 with $p = 7.56\%$ (100,000 replacements made).

⁴Algorithm ?? with large p produces more layer-skipping connections than algorithm 1 with large n ; it can be shown that for any n , algorithm 1 is equivalent to algorithm ?? with $p = N_l - N_l \sqrt{\frac{N_l}{N_o + N_l}} + \frac{N_o}{N_o + N_l} (1 - \frac{N_o + N_l}{N_o N_l})^n$ where N_l is the number of unconnected pairs of neurons that are candidates to share a layer-skipping connection.




Figure 3: Comparison of performance of network topologies on MNIST dataset. Shown are test error (dashed lines) and training error (solid lines). In green is a network with the layered topology and a single learning rate (section 3.2). In orange is a network with the layered topology and unique learning rates (section 3.1), tuned to counter the vanishing gradient problem; this is a recreation of the 5-layer network in [18]. In black is a network with our topology, $p = 7.56\%$ (section 3.3). The network with a layered topology and a single learning rate performs poorly because it suffers from the vanishing gradient problem. The problem can be solved by introducing unique learning rates, or by implementing our topology.

4.1 Network performance comparison

We tracked the error rates of the three networks as they were trained and found that the network with our topology significantly outperforms the basic network with one learning rate, and achieves the same error rates as the basic network with unique learning rates after around 25% more epochs. These results are shown in figure 3.

Notice that the layered network with unique learning rates converges to 0% training error and 2.5% test error in around 150 epochs and the network with our topology converges to 0% training error and 2.5% test error in around 190 epochs, whereas the layered network with one learning rate fails to converge in 250 epochs and has training and test error around .5% higher than the other two networks. While it is possible that the basic network with one learning rate would converge given enough time, it is clearly inferior to the other two networks. It is also apparent that on the MNIST dataset, our network with a single learning rate is practically interchangeable with the layered network with unique learning rates.

4.2 Training rates of individual pairs of layers

We tracked the training rates of individual pairs of layers as described in equations 2 and 3 and found evidence that the extent of the vanishing gradient problem is the primary cause of the performance disparity seen in section 4.1. These results are shown in figure 4.

Notice that in the layered network with one learning rate there




Figure 4: Extent of the vanishing gradient problem in different network topologies while training on MNIST dataset. (left) A network with the layered topology and a single learning rate. The training rate of a layer decreases by a factor of ~ 4 with each additional layer. **(center)** A network with the layered topology and unique per-layer learning rates. The vanishing gradient is counteracted by hand-tuning a learning rate hyperparameter for each layer. **(right)** A network with our topology, $p = 7.56\%$. The new topology eliminates the need for hand-tuning per-layer learning rate hyperparameters.

is a significant spread in the training rates of pairs of layers, with the deepest pair training at around 1% of the rate of the shallowest pair. This problem is solved very effectively in the layered network with unique learning rates. The problem appears to be solved effectively for the deepest 3 pairs in the network with our topology, but the output layer still trains significantly faster than the deeper 3 layers. This makes sense if we assume that the important factor in a layer’s training rate is its expected path length to the target layer, because every neuron in the output layer connects to the target layer through a single connection, whereas paths starting at deeper neurons must first pass through the output layer before connecting to the target layer.

4.3 Error rate after one epoch as connections are added

We tracked the training error after one epoch of a network with our topology while varying p ; the results are shown in figure 5.

Three regimes are evident in the figure. For $p < .1\%$, there is little improvement in the error rate as p is increased, but there is substantial improvement in the uniformity of the training rates of deep layers. When $p > .1\%$, the deep layers are very uniform, and the error rate starts decreasing with p at a rate that is slightly slower than exponential; at this point there is little improvement in the uniformity of deep layers, but the rate of the shallowest layer appears to move closer to those of the deeper layers.

We found that our topology performs significantly worse than the basic topology with one learning rate when few connections are

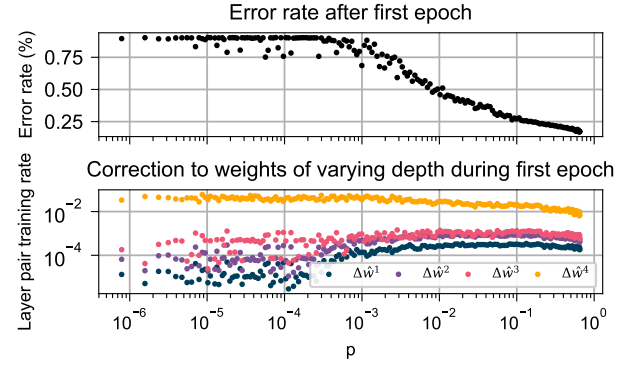


Figure 5: Performance of a network with varying p during the initial stages of training. (top) The training error after one epoch. **(bottom)** Training rate for each layer after one epoch. There is little improvement for $p < .1\%$, after which the error rate decreases at a rate that is slightly slower than exponential with p . Improvement begins after the training rates of non-output layers converge, and improvement beyond that point is associated with convergence between the rates of the output layer and the non-output layers.

replaced. This could be due to a poor weight initialization scheme for the added intralayer connections; we have noticed anecdotally that networks appear to be less-sensitive to their weight initialization scheme as connections are replaced. We have found that networks perform poorly relative to a basic network with one learning rate until p is in the ballpark of 7%. This experiment suggests that training rate will keep improving long after that, but does not show long-term performance or test performance; we suspect that a network’s generalization ability will suffer for large p as it loses its regimented nature.

5 RELATED WORK

Much research has been done in pursuit of a biologically-plausible deep learning algorithm. References [13, 17, 26] describe other algorithms that approximate the gradient of a cost function locally. Reference [14] demonstrates that networks can train through back-propagation with randomly-initialized feedback weights between neurons; the need for symmetric connection weights is not solved by equilibrium propagation and is not biologically-plausible, so this finding could potentially be applied here as well. Reference [4] expands on this work to show that the random feedback matrix can be made very sparse with little loss of network accuracy, significantly reducing need for data transfer within a large network that would otherwise grow quadratically with network size. Reference [16] implements equilibrium propagation using spiking neurons; this type of activation function is believed to be more biologically-plausible than a hardened sigmoid function. Reference [2] discusses the criteria a biologically-plausible algorithm would need to satisfy. Reference [1] surveys promising biologically-motivated algorithms and evaluates their effectiveness on hard datasets.

References [5, 15, 20] discuss neuromorphic architecture that

could potentially implement equilibrium propagation.

Layer-skipping connections have been explored in other contexts. References [7, 23] use layer-skipping connections as a linear transformation in parallel with nonlinear layers to great effect in very-deep convolutional networks. References [11, 25] apply small-world topology to networks trained using backpropagation.

References [6, 10] effectively counter the vanishing gradient problem in the context of deep networks trained using backpropagation.

6 DISCUSSION AND FUTURE RESEARCH

6.1 Characteristics of an effective network

Our experiments indicate that networks with our topology perform poorly relative to a layered network with a single learning rate until p is in the ballpark of 8%. It seems likely that the training rate will continue to improve as p is increased to around 70%; however, we suspect this improvement would be at the cost of a higher test error rate, because as p increases a network's topology looks less-layered and more like a sparsely-connected single-layer network. We did not observe deterioration of a network's generalization ability for $p = 7.56\%$; it is possible that such deterioration is not evident when training on a dataset as simple as MNIST, but would become apparent on a more-difficult dataset such as CIFAR or ImageNet.

6.2 Factors influencing network performance

Contrary to our expectations, we found that a network's performance did not track closely with any of the small-world metrics (Appendix ??). Good performance was achieved when algorithm 1 was executed with $p \approx 8\%$, which leads to a network with small L and C that was therefore not small-world. In retrospect we believe these metrics to be poor predictors of network performance. One reason for this may be that the characteristic path length L quantifies the shortest path between two neurons, but the total number of paths is also important because a neuron's behavior is dictated by an affine transformation of all its inputs. Similarly, it can be important for a network to retain its layered nature in order to generalize well, but the clustering coefficient C does not appear to be a good metric for this - e.g. a layered network without intralayer connections has a small clustering coefficient because the neighbors of a given neuron are not connected to one-another.

There are several factors other than path length that we believe affect performance with respect to p , and that may disguise relationships between performance and the small-world metrics. It has been shown [7, 10] that in deep convolutional networks, performance can be improved by using layer-skipping connections to create an identity mapping or dimension-preserving linear transformation in parallel with groups of nonlinear layers to reframe their task as learning the residual of the output of earlier layers; in our implementation layer-skipping connections form a linear transformation of earlier layers, and this fact may improve performance even though only some of the dimensions of the deeper layers are transferred. Due to the discrete approximation of the network's dynamics whereby a neuron l layers from the output is not influenced by a perturbation of the output until iteration $l + 1$ of the approximation, deep layers experience less evolution

than shallow layers during the weakly-clamped phase; since layer-skipping connections allow deeper layers to begin evolving sooner, they probably reduce this effect. As connections are replaced, a network looks less layered and more like a sparsely-connected single-layer network, so we expect that there is a tradeoff between the rate at which a network trains and its ability to generalize. In this paper our topology did not deteriorate test error relative to the original layered network, but it is possible that a network on a harder dataset such as CIFAR or ImageNet or with a higher value of p would see deterioration as its topology deviates from a layered topology.

6.3 Directions for Future Research

There are several directions in which future research could be taken:

- Finding and mathematically-justifying a weight initialization scheme for added and replaced connections.
- Exploring the effect of adding (instead of replacing) layer-skipping connections, training a network, then removing the connections and training further, to see if there is a lasting benefit after removal; if so, connections could be used to quickly train a network in analog hardware, then removed to reduce the network's power consumption and heat load.
- Exploring the effect of allowing non-output connections to be directly affected in the weakly-clamped phase of training, in effect allowing neurons outside the output layer to connect to the target layer; this would likely reduce the gap between the training rates of output and hidden layers seen in figure 4.
- Exploring a network's test error as p is increased, especially on difficult datasets such as CIFAR or ImageNet where generalization is very important.

The U.S. Government is authorized to reproduce and distribute reprints for governmental purposes notwithstanding any copyright annotation thereon.

REFERENCES

- [1] Sergey Bartunov, Adam Santoro, Blake A. Richards, Geoffrey E. Hinton, and Timothy P. Lillicrap. Assessing the scalability of biologically-motivated deep learning algorithms and architectures. *CoRR*, abs/1807.04587, 2018.
- [2] Yoshua Bengio, Dong-Hyun Lee, Jörg Bornschein, and Zhouhan Lin. Towards biologically plausible deep learning. *CoRR*, abs/1502.04156, 2015.
- [3] E. Bullmore and O. Sporns. Complex brain networks: graph theoretical analysis of structural and functional systems. *Nature*, 2009.
- [4] Brian Crafton, Abhinav Parihar, Evan Gebhardt, and Arijit Raychowdhury. Direct feedback alignment with sparse connections for local learning. *CoRR*, abs/1903.02083, 2019.
- [5] Mike Davies, Narayan Srinivasa, Tsung-Han Lin, Gautham Chinya, Prasad Joshi, Andrew Lines, Andreas Wild, and Hong Wang. Loihi: A neuromorphic manycore processor with on-chip learning. *IEEE Micro*, PP:1–1, 01 2018.
- [6] Xavier Glorot and Yoshua Bengio. Understanding the difficulty of training deep feedforward neural networks. In Yee Whye Teh and Mike Titterton, editors, *Proceedings of the Thirteenth International Conference on Artificial Intelligence and Statistics*, volume 9 of *Proceedings of Machine Learning Research*, pages 249–256, Chia Laguna Resort, Sardinia, Italy, 13–15 May 2010. PMLR.
- [7] Kaiming He, Xiangyu Zhang, Shaoqing Ren, and Jian Sun. Deep residual learning for image recognition. *CoRR*, abs/1512.03385, 2015.
- [8] John Hopfield. Neurons with graded response have collective computational properties like those of two-state neurons. *Proceedings of the National Academy of Sciences of the United States of America*, 81:3088–92, 06 1984.
- [9] Mark Humphries and Kevin Gurney. Network 'small-world-ness': A quantitative method for determining canonical network equivalence. *PLoS one*, 3:e0002051, 02 2008.

- [10] Sergey Ioffe and Christian Szegedy. Batch normalization: Accelerating deep network training by reducing internal covariate shift. *CoRR*, abs/1502.03167, 2015.
- [11] Gokul Krishnan, Xiaocong Du, and Yu Cao. Structural pruning in deep neural networks: A small-world approach, 2019.
- [12] Y. LeCun and C. Cortes. The mnist database of handwritten digits, 1998.
- [13] Dong-Hyun Lee, Saizheng Zhang, Asja Fischer, and Y. Bengio. Difference target propagation. pages 498–515, 08 2015.
- [14] Timothy P. Lillicrap, Daniel Cownden, Douglas B. Tweed, and Colin J. Akerman. Random feedback weights support learning in deep neural networks, 2014.
- [15] Mitchell Nahmias, Bhavin Shastri, A.N. Tait, and P.R. Prucnal. A leaky integrate-and-fire laser neuron for ultrafast cognitive computing. *Selected Topics in Quantum Electronics, IEEE Journal of*, 19:1–12, 09 2013.
- [16] Peter O'Connor, Efstratios Gavves, and Max Welling. Training a network of spiking neurons with equilibrium propagation. 2018.
- [17] Fernando J Pineda. Generalization of back-propagation to recurrent neural networks. *Physical Review Letters*, 59(19):2229–2232, 1987.
- [18] Benjamin Scellier and Yoshua Bengio. Equilibrium propagation: Bridging the gap between energy-based models and backpropagation, 2016.
- [19] Jürgen Schmidhuber. Deep learning in neural networks: An overview. *Neural Networks*, 61:85–117, Jan 2015.
- [20] Jeffrey M. Shainline, Sonia M. Buckley, Adam N. McCaughan, Jeffrey T. Chiles, Amir Jafari Salim, Manuel Castellanos-Beltran, Christine A. Donnelly, Michael L. Schneider, Richard P. Mirin, and Sae Woo Nam. Superconducting optoelectronic loop neurons. *Journal of Applied Physics*, 126(4):044902, Jul 2019.
- [21] Karen Simonyan and Andrew Zisserman. Very deep convolutional networks for large-scale image recognition, 2014.
- [22] Rupesh Kumar Srivastava, Klaus Greff, and Jürgen Schmidhuber. Training very deep networks, 2015.
- [23] Rupesh Kumar Srivastava, Klaus Greff, and Jürgen Schmidhuber. Highway networks. *CoRR*, abs/1505.00387, 2015.
- [24] D. Watts and S. Strogatz. Collective dynamics of 'small-world' networks. *Nature*, 1998.
- [25] L. Xiaohu, L. Xiaoling, Z. Jinhua, Z. Yulin, and L. Maolin. A new multilayer feedforward small-world neural network with its performances on function approximation. In *2011 IEEE International Conference on Computer Science and Automation Engineering*, 2011.
- [26] Xiaohui Xie and Hyunjune Seung. Equivalence of backpropagation and contrastive hebbian learning in a layered network. *Neural computation*, 15:441–54, 03 2003.